

Tags: Python, SQL, revenue, budget, costs, Spark SQL, PostgreSQL, SQL Server, dbt, data build tools, CapEx, OpEx, EC2, AWS,

Python & SQL – Usage Costs, the Good & Wasted

Beginning.....	2
Overall up front:.....	2
Python vs. SQL – Cost Advantages.....	3
Python vs. SQL – Cost Wastes.....	6
A Decent Cost Comparison Table...?.....	8
Check this out on SQL – CTEs / Temp Tables / CTAS.....	9
Temp View.....	9
Temp Tables — Physical Snapshots.....	9
Summary.....	11
Definitions.....	12

Beginning

Sooooo. Using Python and SQL and the costs are adding up eh...? Basically, we need to know about cost models for each area.

Well, for the inexperienced, here is a 'bit' of content to keep in mind to avoid racking up major unnecessary costs.

Much of the material in this paper is basically research I've done to see where there might be better options in different situations. The research is primarily from multiple articles out there but I will also check out a few of the AI chatters to see what they feed back to me.

Over time, I will add more research / experience content here for other new Data Engineers (or ML Engineers).

But, if'n youse seez sumthin' that is wrong or incorrect – feel free to lemme know and explain why youse believe it to be so...

(*One must have humor from time to time; it makes work transpire much easier...* and I have been in MUCH more serious work situations while a Marine a long time ago and more recently while I was working as a civilian/fed intelligence officer.)

Overall up front:

- **SQL** is less expensive for moving, filtering, and aggregating data
- **Python** is less expensive for orchestration, complex logic, and non-relational work
- The **biggest waste** is using Python to do what SQL engines **already does less expensively**

Python vs. SQL – Cost Advantages

Aspect	Python	SQL
Costs, Costs, Costs....	<p>Complex logic/transformations on small-to-medium data</p> <ul style="list-style-type: none"> • No query costs if data is already local • Reusable code without re-querying • Free compute on your local machine <p>Iterative development/prototyping</p> <ul style="list-style-type: none"> • No per-query costs like cloud DBs • Can manipulate data in-memory repeatedly <p>Where Python Has the BIGGEST Cost Advantage</p> <p>Orchestration & glue logic</p> <p>Python is cost effective for:</p> <ul style="list-style-type: none"> • Control flow • Conditional logic • Pipelines • API calls • Multi-system integration <p>Cost:</p> <p>For large datasets:</p> <ul style="list-style-type: none"> • Python = dollars to hundreds of dollars 	<p>Working w/ large datasets in DBs (database)</p> <ul style="list-style-type: none"> • SQL runs on the DB server (close to data) – i.e.: SQL engines (Redshift, Athena, BigQuery, Snowflake) • No data transfer costs - processing happens where <i>that</i> data lives – less network i/o • Highly optimized for filtering / aggregating millions of rows – i.e.: <ul style="list-style-type: none"> ◦ Filtering 1 billion rows in Athena costs pennies vs loading into Python <p>Cost:</p> <ul style="list-style-type: none"> • Scans only required columns • Executes in compiled engine • Parallelized automatically <p>Less expensive than pulling data into Python.</p> <p>For large datasets:</p> <ul style="list-style-type: none"> • SQL = pennies <i>(Still, it will depend on task...)</i> <p>Serverless SQL is brutally cheap, i.e.:</p> <ul style="list-style-type: none"> • Athena • BigQuery on-demand <p>Cloud data warehouse queries (like in Athena usage)</p> <ul style="list-style-type: none"> • Only pay for data scanned, i.e.: using: • TABLESAMPLE SYSTEM (10) in SQL = 90% cost savings immediately <p>Efficient WHERE clauses reduce scanned data</p> <p>Or using specific column selection:</p> <p>Example:</p> <pre>SELECT class, AVG(sepal_length) FROM iris GROUP BY class;</pre>

Aspect	Python	SQL
Costs, Costs, Costs (cont.) ... Flexibility	<p>Small to medium data transformations</p> <p>If data:</p> <ul style="list-style-type: none"> • Fits in memory • Is already filtered • Needs complex logic <p>Python (pandas, numpy) is cheaper than spinning up:</p> <ul style="list-style-type: none"> • Spark Warehouses • Distributed SQL <p>Non-tabular data</p> <p>Python wins for:</p> <ul style="list-style-type: none"> • JSON XML Images • NetCDF Logs ML features <p>It is very versatile and handles various data types & formats; allowing for broader implementation in data science, automation & web applications.</p> <p>Scales linearly with data size</p> <p>Python:</p> <ul style="list-style-type: none"> • Interpreted • GIL-bound • Memory-heavy 	<p>Cost is usually:</p> <ul style="list-style-type: none"> • per TB scanned • No idle compute cost <p>If our query scans 5 GB: You pay for 5 GB — not for a running server.</p> <p>It is specialized for relational DB operations, making it efficient for data manipulation & retrieval.</p> <p>Scales linearly with data size</p> <p>SQL engines:</p> <ul style="list-style-type: none"> • Are written in C/C++ • Vectorized • Highly optimized

Aspect	Python	SQL
Flexibility	<p>It is very versatile and handles various data types & formats; allowing for broader implementation in data science, automation & web applications.</p> <p>Scales linearly with data size</p> <p>Python:</p> <ul style="list-style-type: none"> • Interpreted • GIL-bound • Memory-heavy 	<p>It is specialized for relational DB operations, making it efficient for data manipulation & retrieval.</p> <p>Scales linearly with data size</p> <p>SQL engines:</p> <ul style="list-style-type: none"> • Are written in C/C++ • Vectorized • Highly optimized
Development Costs	<p>It often requires fewer lines of code for complex tasks due to its readability and powerful libraries (e.g., Pandas, NumPy), which can reduce development time.</p>	<p>It can be more straightforward for querying DBs, which can lead to faster execution of specific data retrieval tasks without extensive coding.</p>
Community & Libraries	<p>A large array of libraries & frameworks is available for numerous tasks, reducing the need for custom development.</p>	<p>SQL's standardized queries are supported across various DB systems, lowering training costs for teams familiar w/ SQL syntax.</p>
Learning Curve	<p>Although Python has a learning curve, it is generally considered easier for beginners, which reduces training costs.</p>	<p>SQL is straightforward for data querying, but advanced features and optimizations may require deeper expertise, potentially increasing training expenses.</p>

Python vs. SQL – Cost Wastes

Python	SQL
<ul style="list-style-type: none">Performance Overhead: Python may experience slower performance for large datasets compared to SQL. And if you do not optimize it, it will lead to higher computational costs.Library Management: Over-reliance on multiple libraries can lead to version conflicts and increased maintenance costs. A major headache for everyone...Inefficient Code: Poorly written Python code leads to longer execution times, demanding more computational resources and increasing operational (OpEx) costs. Possibly, OpEx budget might balloon but 'could' still be better than CapEx costs...	<ul style="list-style-type: none">Licensing Costs: Some SQL DBMS (<i>database management systems</i>) may involve licensing fees, leading to higher OpEx costs compared to open-source alternatives.Complexity in Scaling: As DBs grow, scaling SQL solutions can become expensive and complex, especially if additional features or optimizations are needed.Training and Expertise: Advanced SQL queries and optimizations may require specialized knowledge, leading to higher labor costs for skilled personnel.<ul style="list-style-type: none">Here, this could add to OpEx budget more than expected...
Loading massive datasets into memory # EXPENSIVE: Pulls 10 GB into memory/network df = pd.read_sql("SELECT * FROM big_ole_table", connection)	One that everyone should already know about – the use of the infamous '*' Full table scans without filters -- Expensive: Scans entire table SELECT * FROM big_ole_table;
# BETTER: Filter in SQL first df = pd.read_sql("SELECT * FROM big_ole_table WHERE ...", connection)	-- Less expensive: Only scans needed data SELECT * FROM big_ole_table WHERE date = '2026-01-10';
Data transfer costs <ul style="list-style-type: none">Moving data from cloud → local machine = bandwidth costsProcessing 100 GB locally vs in-DB = huge transfer fees	Not using partitioning in cloud databases <ul style="list-style-type: none">Athena charges by data scannedQuerying unpartitioned tables = scanning everything
Inefficient loops on large data # Slow & Expensive (CPU time) for row in df.iterrows(): # Don't do this! process(row)	Repeated identical queries <ul style="list-style-type: none">Each query costs money in cloud DBsShould cache results or use materialized views
# FAST & Less expensive df['result'] = df.apply(vectorized_function) # Vectorized operations	SELECT * when you only need a few columns -- Wastes money scanning unused columns SELECT * FROM table; -- Better SELECT col1, col2 FROM table;

Python

SQL

Not using appropriate compute resources

- Running big jobs on expensive EC2 instances when Lambda / Glue would be cheaper
- Keeping instances running 24/7 for occasional processing ()

(This is an area at one time, where I was running point at an intel agency. To have all divisions use AWS cost explorer and Bill & Management – to review and revise their use of EC2s [*many had EC2s that were complete overkill for the task at hand, especially for DevOps & Testing*], **DBs and S3 storage** [*many had 100s of TBs of data needlessly in Standard – not being accessed*]).

Few months later, final review: **\$100s of Ks / month saved...**

Biggest Cost WASTE with Python

Pulling large datasets out of DBs -- **Classic mistake:**

```
df = pd.read_sql("SELECT * FROM  
big_ole_table", conn)  
df = df.groupby("class").mean()
```

WHY is it expensive...?

- Full table scan Network transfer
- Python memory usage Longer runtimes
- Bigger instances

FIX:

```
SELECT class, AVG(sepal_length)  
FROM huge_table  
GROUP BY class;
```

Cost-effective workflow:

- 1) **SQL:** Filter/aggregate/sample in Athena (cheap, close to data)
- 2) Export: Only move reduced dataset
- 3) **Python:** Complex visualizations/analysis locally (free)

i.e.:

```
SELECT lat, lon, temp_c, forecast_step  
FROM mogreps_analytics.mogreps_latest_6h  
TABLESAMPLE SYSTEM (10)    # 90% cost  
savings!  
WHERE forecast_step < 6      # Further  
reduces scanned data
```

SQL is **terrible** at:

- 1) Loops
- 2) Branching
- 3) External APIs
- 4) Complex business logic

Trying to do 1 - 4 in SQL leads to:

- Stored procedure abuse
- Long-running warehouse sessions
- Developer inefficiency (which is a real cost)

SQL engines either:

- Scan too much
- Require expensive UDFs
- Force schema contortions

Python	SQL
Using Python where vectorized SQL exists	Using SQL warehouses as application logic engines:
(a) Examples: <ul style="list-style-type: none"> • Filtering • Aggregations & Joins • Deduplication (c) Overusing Spark/Glue for small jobs Running Glue for: <ul style="list-style-type: none"> • 10 MB .csv • Simple aggregation 	(b) Python does these: <ul style="list-style-type: none"> • Slower • More memory • More compute cost (d) Cost waste: <ul style="list-style-type: none"> • Job spin-up • Executor overhead • Logging overhead Lambda + SQL query would be cheaper
	(a) Running: <ul style="list-style-type: none"> • Long sessions • Stored procedures w/ loops • Repeated transformations (c) Re-scanning massive tables repeatedly <pre>SELECT ... FROM fact_table WHERE date = '2024-01-01';</pre> Run 100 times: <ul style="list-style-type: none"> • We pay for 100 scans
Cost wasteful workflow: <ol style="list-style-type: none"> 1) Worst → SELECT * 2) Pull entire table into Python or <ul style="list-style-type: none"> • pulling raw data instead of aggregated data 3) Filter/sample in pandas 4) Pay for data transfer AND compute 	(a) Doing ML or complex math in SQL – bad... <ul style="list-style-type: none"> • Iterative algorithms • ML training • Feature engineering with state (b) Result: <ul style="list-style-type: none"> • Long-running queries • High warehouse costs • Poor maintainability Python is cheaper here.
Do control logic in Python Python minimizes developer and orchestration cost. As some say, " Let Python orchestrate — not compute. " AND NEVER... Never move large data just to "feel comfortable."	Do data reduction in SQL SQL minimizes compute cost per byte.

A Decent Cost Comparison Table...?

Task	Cheaper Tool	Why
Filtering rows	SQL	Predicate pushdown
Aggregations	SQL	Vectorized execution
Joins	SQL	Optimized planners
API calls	Python	SQL can't
Complex business rules	Python	Readability + speed
ETL orchestration	Python	Control flow
ML training	Python	Libraries + GPUs
BI datasets	SQL	Columnar + caching
Small ad-hoc scripts	Python	No warehouse spin-up

Check this out on SQL – CTEs / Temp Tables / CTAS

CTEs (common table expression) – every single time that CTE is run (*in the overall use of one application*), we are:

NOTE: The CTE is a **query rewrite** AND a **logical alias**

- a) creating something that is run every time it is called,
- b) *s-l-o-w-i-n-g* down the application processing time (*resource usage – data shuffling & I/O costs*),
- c) every time it is called in the code, we undergo computational costs, needlessly,
 - **ESPECIALLY** when they use a **SELECT *** statement...
 - But. The SQL CTE is not inherently ‘always’ creating more computation than a Temp View
- d) The CTE is scoped to one query and disappears after execution
- e) And as one person wrote in a comment (*Portuguys on Medium*) – “it’s not a binary choice: never / always use CTEs – the scenario is, it depends.”

```
WITH my_cte AS (
  SELECT * FROM large_table WHERE condition
)
SELECT * FROM my_cte;
```

Engines that materialize CTEs by default	
Some systems (historically): <ul style="list-style-type: none"> • Older PostgreSQL • Some MPP engines • Certain warehouse configurations 	They treat CTEs as optimization fences Result: <ul style="list-style-type: none"> • Forced materialization • Extra I/O

Temp View

- **Session-scoped** - persists across multiple queries
- **Not materialized** (if it's a VIEW) - also recomputed each time
- **Reusable** - can query it multiple times in different statements
- **But – still not stored unless** you use:
 - CREATE TEMP TABLE AS

```
CREATE TEMP VIEW my_view AS
SELECT * FROM large_table WHERE condition;

SELECT * FROM my_view;
```

Caching **Temp views/tables** avoids repeated scans/shuffles when performance is a priority.

The catch, thinking this way may be more applicable for **PostgreSQL** more than for **SQL Server**.

But. CTEs and Temp Views are about the same. Neither of the two are materialized (*not physically stored*), so any computation occurs when a query is called.

What CTEs have going for it is that they **are** very readable for clear thought reasoning (*easy explanation*)...

Temp Tables — Physical Snapshots

Sometimes, you need intermediate results to *persist physically* for debugging or reuse across sessions.

That's when **TEMP tables** shine.

Materialized vs Non-Materialized			CTEs vs	TEMP Tables/CTAS
Type	Materialized?	Computation	Use CTEs when:	Use Temp Tables/CTAS when:
CTE	No	Computed every reference	<ul style="list-style-type: none"> • referencing result once or twice • query is simple and fast • want clean, readable code 	<ul style="list-style-type: none"> • querying result many times • computation is expensive • working w/ large datasets • Incremental pipelines
Temp VIEW	No	Computed every reference		
Temp TABLE	Yes	Computed once , then stored		

Use CTEs when: <ul style="list-style-type: none"> • We reference the result once or twice • The query is simple & fast • We want clean, readable code 	Use Temp Tables/CTAS when: <ul style="list-style-type: none"> • We query the result many times • The computation is expensive • working w/ large datasets
---	--

CTE Uses MORE Computation:	TEMP Tables – a better approach:
<p>-- BAD: CTE referenced 3 times = computed 3 times</p> <pre>WITH expensive_cte AS (SELECT complex_join(...) FROM massive_table) SELECT * FROM expensive_cte UNION ALL SELECT * FROM expensive_cte -- Recomputed! UNION ALL SELECT * FROM expensive_cte; -- Recomputed again!</pre>	<p>-- GOOD: Materialized once</p> <pre>CREATE TEMP TABLE expensive_result AS SELECT complex_join(...) FROM massive_table; SELECT * FROM expensive_result UNION ALL SELECT * FROM expensive_result -- Just reads stored data UNION ALL SELECT * FROM expensive_result;</pre>

Expensive Cost implication:	Better approach efficient in Athena:
<p>-- If you reference a CTE multiple times:</p> <pre>WITH sampled_data AS (SELECT * FROM mogreps_latest_6h #a View of mine TABLESAMPLE SYSTEM (10)) SELECT * FROM sampled_data WHERE condition1 UNION ALL SELECT * FROM sampled_data WHERE condition2;</pre> <p>This might scan the base table TWICE (depending on query optimizer)</p>	<p>-- Create a real table (one-time scan cost)</p> <pre>CREATE TABLE mogreps_sample AS ---- # a table of mine SELECT * FROM mogreps_latest_6h---- # a View of mine TABLESAMPLE SYSTEM (10);</pre> <p>-- Now query the smaller table multiple times (cheap!)</p> <pre>SELECT * FROM mogreps_sample WHERE condition1; SELECT * FROM mogreps_sample WHERE condition2;</pre>

In using AWS Athena	
Athena does NOT support	Athena DOES support
<ul style="list-style-type: none"> • Temp tables • Temp views • Materialized views 	<ul style="list-style-type: none"> • CTEs (WITH clause) • Creating real tables with <ul style="list-style-type: none"> ◦ CREATE TABLE AS SELECT

Cost comparison summary		
Scenario	Cheaper	Why
Single-use logic	CTE	Inlined
Simple filtering	Either	Same plan
Reused many times	Temp table	One compute
Large expensive joins	Temp table	Avoid recompute
Glue / Spark	Same	Logical rewrite
Warehouses w/ materialized CTEs	Temp table	Predictable cost

Making good use of **dbt** (*data build tools*) is a good SQL tool, though I, myself, have not had much use of it yet. I am going by a couple of folks who make the tool credible.

But I do know that we have to be careful with **dbt**, regarding the introduction of additional dependencies.

Temporary views from **Spark SQL** can be connected to caching to avoid repeated computation of heavy intermediate logic.

Some Rules of Thumb:

- **"Push computation to the data"** (use SQL where data lives)
- **"Pull only what you need"** (filter/sample before extraction)
- **Use Python for logic SQL can't handle** (ML, complex transformations, visualizations)

For a Data Engineer project I did, I used AWS Lambda + Glue (ELT) + S3+ Athena. Here, I chose to do reduced sampling (10%) instead of using the entire dataset. In this case, a Glue Python Shell job was better than the Spark Job. The source dataset was only ~22 MB but it was in .nc format, which when expanded, balloons out into a **6.6 GB dataset** (.csv). So, 10% sampling in this scenario...

For visualization, after the ELT job, the 10% sampling in SQL is more cost-effective than pulling everything and sampling in Python.

Summary

Python might offer more: flexibility, ease of use, and a broad range of applications. Whereas, SQL excels in structured data retrieval and manipulation.

Each tool comes with its very own cost advantages and potential waste, making the company's choice of which, very dependent on project-specific requirements and long-term resource management strategies.

For every task, careful consideration and expected scalability can help mitigate unnecessary CapEx AND OpEx costs in both Python and SQL environments.

Definitions

CapEx: capital expenses, traditional – one-time purchases:

- buying servers/hardware; building a data center; purchasing software licenses (perpetual); buying a company car – basically, large upfront costs, that depreciate over time with values also decreasing...

OpEx: operating expenses – to run business day-to-day:

- rental/leasing costs for infrastructure, equipment insurance, site lease/rent, cloud services (AWS, Azure), i.e.: Glue/Athena query costs, monthly software subscriptions (Office 365, Salesforce), personnel (salaries), utilities (heating/cooling/power), internet...
- regular recurring payments; fully deductible in the year spent; no ownership – we're renting/subscribing; scales with usage (pay as you go)